

image_classification_with_neural_networks

December 23, 2025

1 Classification with Neural Networks

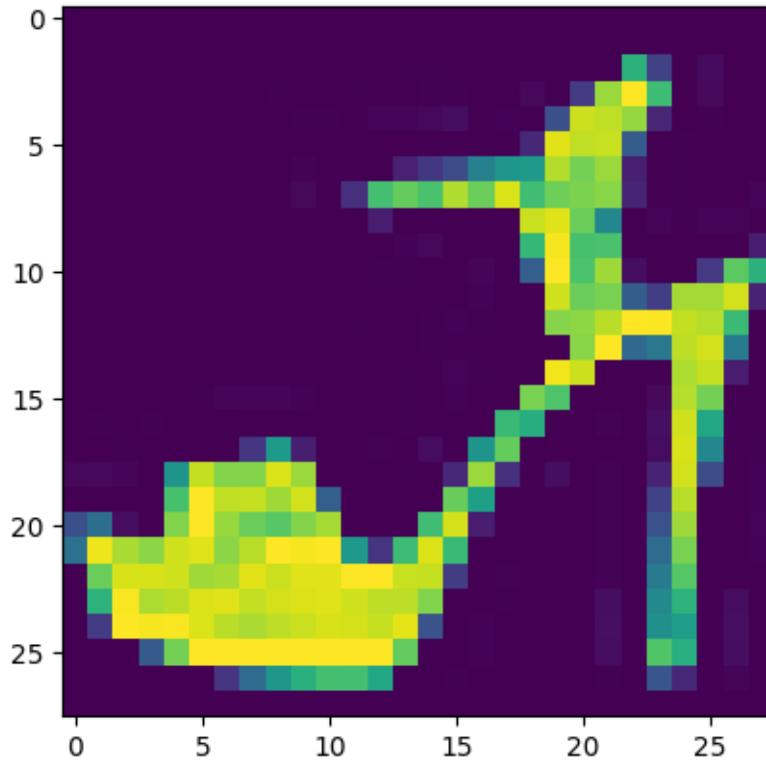
```
[1]: import tensorflow as tf
from tensorflow import keras
import numpy as np
import matplotlib.pyplot as plt
fashion = keras.datasets.fashion_mnist
(xtrain, ytrain), (xtest, ytest) = fashion.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
29515/29515          0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26421880/26421880      4s
Ous/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
5148/5148            0s 3us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4422102/4422102       2s
Ous/step
```

```
[2]: imgIndex = 9
image = xtrain[imgIndex]
print("Image Label : ",ytrain[imgIndex])
plt.imshow(image)
```

Image Label : 5

```
[2]: <matplotlib.image.AxesImage at 0x25700cc7750>
```



```
[3]: print(xtrain.shape)
      print(xtest.shape)
```

```
(60000, 28, 28)
(10000, 28, 28)
```

```
[4]: model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(300, activation="relu"),
    keras.layers.Dense(100, activation="relu"),
    keras.layers.Dense(10, activation="relu"),
])
print(model.summary())
```

```
C:\Users\USER\miniconda3\envs\ds4b\Lib\site-
packages\keras\src\layers\reshaping\flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)

Model: "sequential"
```

| Layer (type) | Output Shape | |
|-------------------|--------------|---|
| Param # | | |
| flatten (Flatten) | (None, 784) | □ |
| 0 | | |
| dense (Dense) | (None, 300) | □ |
| 235,500 | | |
| dense_1 (Dense) | (None, 100) | □ |
| 30,100 | | |
| dense_2 (Dense) | (None, 10) | □ |
| 1,010 | | |

Total params: 266,610 (1.02 MB)

Trainable params: 266,610 (1.02 MB)

Non-trainable params: 0 (0.00 B)

None

```
[5]: xvalid, xtrain = xtrain[:5000]/255.0, xtrain[:5000]/255.0
      yvalid, ytrain = ytrain[:5000]/255.0, ytrain[:5000]
```

```
[6]: model.compile(loss = "sparse_categorical_crossentropy",
                  optimizer = "sgd",
                  metrics=["accuracy"])
history = model.fit(xtrain, ytrain, epochs=30,
                     validation_data=(xvalid, yvalid))
```

```
Epoch 1/30
157/157          5s 16ms/step -
accuracy: 0.2554 - loss: 3.3712 - val_accuracy: 0.0000e+00 - val_loss: 18.0639
Epoch 2/30
157/157          2s 12ms/step -
accuracy: 0.0976 - loss: 6.8426 - val_accuracy: 0.0000e+00 - val_loss: 18.0639
Epoch 3/30
157/157          2s 12ms/step -
accuracy: 0.0976 - loss: 6.8425 - val_accuracy: 0.0000e+00 - val_loss: 18.0635
Epoch 4/30
157/157          2s 11ms/step -
accuracy: 0.0976 - loss: 8.0733 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 5/30
```

```
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 6/30
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 7/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 8/30
157/157      2s 12ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 9/30
157/157      3s 17ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 10/30
157/157      2s 12ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 11/30
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 12/30
157/157      2s 15ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 13/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 14/30
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 15/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 16/30
157/157      3s 22ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 17/30
157/157      2s 12ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 18/30
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 19/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 20/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 21/30
```

```

157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 22/30
157/157      2s 12ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 23/30
157/157      2s 14ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 24/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 25/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 26/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 27/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 28/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 29/30
157/157      2s 10ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099
Epoch 30/30
157/157      2s 11ms/step -
accuracy: 0.0976 - loss: 8.2680 - val_accuracy: 0.0000e+00 - val_loss: 17.9099

```

```
[7]: new = xtest[:5]
predictions = model.predict(new)
print(predictions)
```

| | 0. | 1744.8864 | 2466.2341 | 37160.547 | 1103.994 |
|-----------|-----------|-----------|-----------|-----------|-----------|
| [[| 3765.871 | 1593.8174 | 0. | 0. |] |
| [| 0. | 6121.2056 | 8705.096 | 127639.47 | 3077.999 |
| 13245.684 | 5246.368 | 0. | 0. |] | |
| [| 0. | 4294.097 | 6181.812 | 92248.51 | 2243.1167 |
| 9485.057 | 3808.6047 | 0. | 0. |] | |
| [| 0. | 2969.3694 | 4282.6733 | 63950.645 | 1558.9495 |
| 6552.636 | 2695.8813 | 0. | 0. |] | |
| [| 0. | 3858.3806 | 5489.216 | 81062.125 | 1965.4791 |
| 8395.744 | 3368.7114 | 0. | 0. |]] | |

```
[ ]:
```